

SAI ANEESH GANTI

✉ gantisaianeesh@gmail.com | 🐙 github.com/Saianeesh2003 | 🔗 linkedin.com/in/sai-aneesh-ganti | 📞 9948668997

SUMMARY

AI Engineer and Backend Developer with hands-on experience building production-ready RAG pipelines, agentic LLM workflows, and full-stack applications. Skilled in designing multi-step AI systems (LangGraph, tool calling, vector search) and backend infrastructure (REST APIs, PostgreSQL, authentication). Focused on shipping reliable, end-to-end systems from AI reasoning layers to the backend services that support them.

TECHNICAL SKILLS

Programming: Python, JavaScript (Node.js), SQL

Backend: Node.js, Express.js, REST API Design, PostgreSQL, Drizzle ORM, MVC Architecture, JWT Authentication

AI & GenAI: Agentic Workflows, RAG Systems, Tool Calling, LLM Integration (OpenAI, Claude), Prompt Orchestration, Vector Databases (Qdrant, Pinecone), LangGraph

Libraries: LangGraph, Pandas, NumPy, Scikit-learn, Streamlit

ML Fundamentals: Classification, Regression, Model Evaluation, Data Preprocessing

Tools: Git, Docker, LangSmith, MongoDB, Postman, Jupyter, Streamlit Cloud

EXPERIENCE

IT Intern – IKCON Digital IT Services Pvt Ltd

June 2026 – Present

Onsite, Hyderabad, India

- Designed and developed Cortex, an AI chatbot for the company's product Unified Management Suite (UMS), from early-stage architecture through implementation
- Built a hybrid retrieval pipeline (dense and sparse search, Reciprocal Rank Fusion, cross-encoder reranking) alongside a deterministic SQL fast-path, cutting latency and hallucination on common queries
- Implemented row-level access control across both retrieval paths so answers respect organizational reporting hierarchy and meeting attendance
- Built a RabbitMQ-based sync and reconciliation pipeline keeping PostgreSQL and the vector index consistent with an external source system
- Designed and developed SURAKSHA GRID, a full-stack emergency coordination platform enabling real-time incident reporting, AI-driven impact prediction, and cross-department coordination across Hyderabad's GHMC, HYDRAA, Traffic Police, and other agencies
- Integrated local LLM's through ollama for AI-driven prediction
- Implemented self-hosted WebRTC video conferencing via LiveKit for real-time cross-department coordination

Gen AI Intern – Cynchrony

Aug 2025 – Mar 2026

Remote

- Built and iterated on end-to-end RAG pipelines for document understanding: ingestion, semantic chunking, embedding generation, vector storage, and retrieval
- Designed multi-step agentic workflows using LangGraph, enabling structured reasoning, dynamic routing, and tool-based execution
- Developed LLM pipelines to transform unstructured inputs (PDFs, transcripts) into structured outputs for downstream analysis
- Implemented prompt orchestration and fallback strategies to reduce hallucinations and improve reliability in LLM responses
- Built tool-calling agents capable of multi-step execution across APIs and external tools
- Performed evaluation and debugging of LLM systems through prompt iteration, retrieval tuning, and response analysis
- Deployed lightweight AI applications using Streamlit, focusing on rapid prototyping and usability

PROJECTS

JEE Math Mentor – Multi-Agent AI System | 🐙 GitHub | Live Demo

- Built a 5-agent LangGraph pipeline (Parser to Router to Solver to Verifier to Explainer) to solve JEE-level math problems end-to-end
- Integrated RAG-based retrieval using Pinecone and embeddings for topic-specific formula access

- Added multimodal capabilities: OCR (EasyOCR) and speech input (Whisper) with human-in-the-loop validation
- Designed memory system using MongoDB to retrieve and reuse past solutions
- Designed agent coordination logic to handle failure cases (incorrect parsing, retrieval mismatch), improving solution reliability across multi-step reasoning
- Instrumented system with LangSmith tracing and cost tracking (\$0.008/request)

AI Agent with Autonomous Tool Calling | GitHub | [Live Demo](#)

- Built an agent capable of dynamic tool selection and multi-step execution across APIs
- Implemented reasoning loop for tool selection, parameter extraction, and task orchestration
- Designed decision logic for tool selection based on task decomposition, enabling adaptive multi-step execution
- Designed fallback mechanisms to handle failures and incomplete outputs
- Delivered production-ready MVP within 2 weeks

RAG-Based Document QA System | GitHub | [Live Demo](#)

- Implemented full pipeline: PDF ingestion to chunking to embeddings to vector storage to retrieval to LLM response
- Used Qdrant for semantic search and context management
- Tuned retrieval strategies (top-k, similarity thresholds) for improved answer relevance
- Designed prompts for context-aware generation with reduced hallucination

URL Shortener – Backend REST API | GitHub

- Built a Node.js/Express REST API for creating and managing shortened URLs, with JWT-based authentication and per-user access control
- Designed the schema and queries with PostgreSQL and Drizzle ORM, including unique short-code generation and collision handling
- Followed REST API design and MVC architecture principles, with request validation via Zod

EDUCATION

Osmania University, Hyderabad

Bachelor of Engineering – Electronics and Communication Engineering

July 2025

GPA: 8.09/10

ADDITIONAL INFORMATION

Strengths: Execution focus – built and deployed multiple production AI systems independently; fast learner – self-taught LangChain, LangGraph, and vector databases from scratch; ownership – comfortable taking systems from idea to build to deploy to iterate

Achievement: Selected for MeitY Chips-to-Startup (C2S) Program

Current Focus: Multi-agent systems, production RAG pipelines, backend development with Node.js